

Thursday, July 18, 2019 6:30 AM

Baseline Freewin Sequential
Scheduled Freewin Sequential
Scheduled WORM Sequential
Baseline Freewin Random
Scheduled Freewin Random
Scheduled WORM Random

Freewin == Baseline WORM, because
be better instead of 8%.

READ ONLY
40% REDUCES

READ ONLY
40% REDUCES

A	Chicks 8
B	Chicks 20
C	Free style Fast
D, etc.	Other ciphers


X	chacha8	x	chacha20
Y	chacha20	x	Freestyle Fast
Z	chacha8	x	Freestyle Fast

↑ primary swap ↑


Ratio			
1	$\frac{1}{4}x$	sup,	$\frac{3}{4}x$ primary
2	$\frac{2}{4}x$	sup,	$\frac{2}{4}x$ primary
3	$\frac{3}{4}x$	sup,	$\frac{1}{4}x$ primary

<u>Scheduler</u>	
M	Microed
S	Selective
F, G, H	Forward 0, 1, 2 (respectively)


1.1 Baseline cipher performance, 40MB reads, no scheduling




Baseline cipher performance, 40MB reads,
no scheduling



2.1 Cipher pair performance vs baseline, 40MB reads, forward 0 scheduling, all ratios



3.1 Cipher pair performance vs baseline, 40MB reads, all forward schedules, all ratios



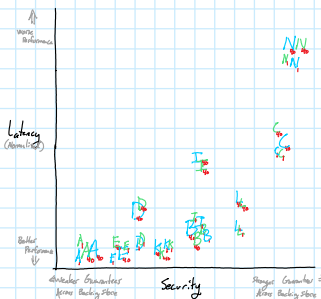
Along w/ figures 4, 7, and 11 from v1

$A \quad B \quad C \quad \left. \vphantom{A \quad B \quad C} \right\} \text{ Bipolar results are simple}$
 $A_R \quad C_W \quad \left. \vphantom{A_R \quad C_W} \right\} \text{ For when read and writes appear on the same graph}$
 $\rightarrow Z_R \quad FZ_R \quad \left. \vphantom{\rightarrow Z_R \quad FZ_R} \right\}$
 $MX_i \leftarrow \text{also other}$

short hand
if & schedule
if not specified


- All examples could exist for more than just 40MB results
- Normalised Y axis to $\min(Y)$ where appropriate
- Expand to multiple charts w/ all combinations for distribution
- Security of cipher pass can be calculated w/ : $S = 4 + 16 \cdot \log_2 \lambda$, $\lambda = \text{ratio}$

1.2 **Baseline cipher performance, 40MB vs 1KB reads, no scheduling** (could separate into two charts instead)




1.2-41t (could separate into two clients instead)

Baseline cipher performance, 40MB vs 1KB reads, no scheduling



2.1-alt


Cipher pair performance vs baseline, 40MB reads, forward 0 scheduling, all ratios



1.3


(could separate into
num clients instead)

Baseline cipher performance, 40MB reads vs writes,
no scheduling



1.3-4H (could separate into two charts instead)

Baseline cipher performance, 40MB reads vs writes, no scheduling



Argues that a rich tradeoff space exists between stream ciphers in a storage context


Demonstrates that we have a (good) mechanism to navigate this tradeoff space

Could also draw best-fit lines to make point?

Would this be a better argument
with the corner for $\text{char}16_t \times \text{char}16_t$
and $\text{char}16_t \times \text{FreeStyle}$ instead of $\text{char}16_t \times \text{FreeStyle}$?

Shows Forward scheduling potential advantages
versus each other and baseline (frontier)

Also figure 7 from v1



with the cipher for chacha8 x freestyle and chacha20 x freestyle. instead of chacha8 x freestyle?

Shows Forward scheduling potential advantages versus each other and baseline (forward)

Also figure 7 from v1


Valid comparison for the different workloads yield similar results?

Good graph for the second use case w/ attempting to shift the storage into a "low power" mode by scheduling cipher jumps from freestyle first to chacha8 as the ratio shifts towards 1/4th of the backing store being freestyle

Also good for the fourth use case w/ SSD End for the same reason: we can dynamically shift into a higher performance mode when the system determines it is necessary but shift back afterwards.

4.1

Cipher partition performance, 40MB reads, mirrored vs selective schedules vs baseline



Also Figure 11 from v1


Good graph for the first and third use cases: VSRs and "pass mode". In the case of VSRs, lack of perf degradation despite presence of VSRs (up to a point). In the case of pass mode, see Figure 11.

I/O is not "ratio"-ed; operations occur on one partition or the other and the results are measured

Shows selective & mirrored are parity w/ baseline (+ overhead)

5.1

Cipher partition performance, 40MB writes, mirrored vs selective schedules vs baseline



Shows that mirrored is slow for writes but as fast as baseline & selective for reads. Figure 11 shows the benefit of using mirrored schedule.